

Real-time Data Integration and Analytics Platform on AWS for IoT Applications

**Proposed Project and Progress Report
Cloud Computing - AT 83.03**

Tanzil Al Sabah - ST123845 - AIT IoT

Background

The Internet of Things (IoT) has dramatically transformed various sectors by enabling the collection and exchange of data from a myriad of connected devices. This transformation is facilitating unprecedented levels of automation, efficiency, and insight across industries such as healthcare, manufacturing, agriculture, and smart cities. As IoT devices proliferate, they generate vast quantities of data that need to be processed and analyzed in real-time to extract actionable insights. This requirement presents significant challenges, primarily due to the volume, velocity, and variety of the data produced.

Cloud computing, with its scalable and flexible resource management capabilities, has emerged as a viable solution to address these challenges. Specifically, AWS provides a comprehensive suite of services that support the development of scalable, efficient, and cost-effective real-time data processing and analytics platforms. However, designing a system that can seamlessly integrate these services to handle the complexities of IoT data in real-time demands a strategic approach and careful consideration of architecture, data flow, and processing logic.

Problem Statement

Despite the potential of IoT devices to revolutionize data-driven decision-making, several challenges hinder the efficient processing and analysis of IoT data in real-time. These challenges include:

Scalability: Traditional data processing solutions often cannot scale dynamically to accommodate the fluctuating volumes of data generated by IoT devices.

Latency: Many IoT applications require low-latency processing and analysis to deliver timely insights, a requirement difficult to meet with conventional data processing architectures.

Data Integration and Management: The heterogeneity of IoT devices and the data they produce complicates integration and management, necessitating flexible data processing pipelines that can adapt to diverse data formats and sources.

Cost: Efficiently managing the cost of cloud resources while scaling to meet demand is a persistent challenge for real-time IoT data processing platforms.

Security and Privacy: Ensuring the security and privacy of data throughout its lifecycle, from ingestion to analytics, is paramount, especially given the sensitive nature of data collected in some IoT applications.

Objectives

To address the challenges identified, this project aims to:

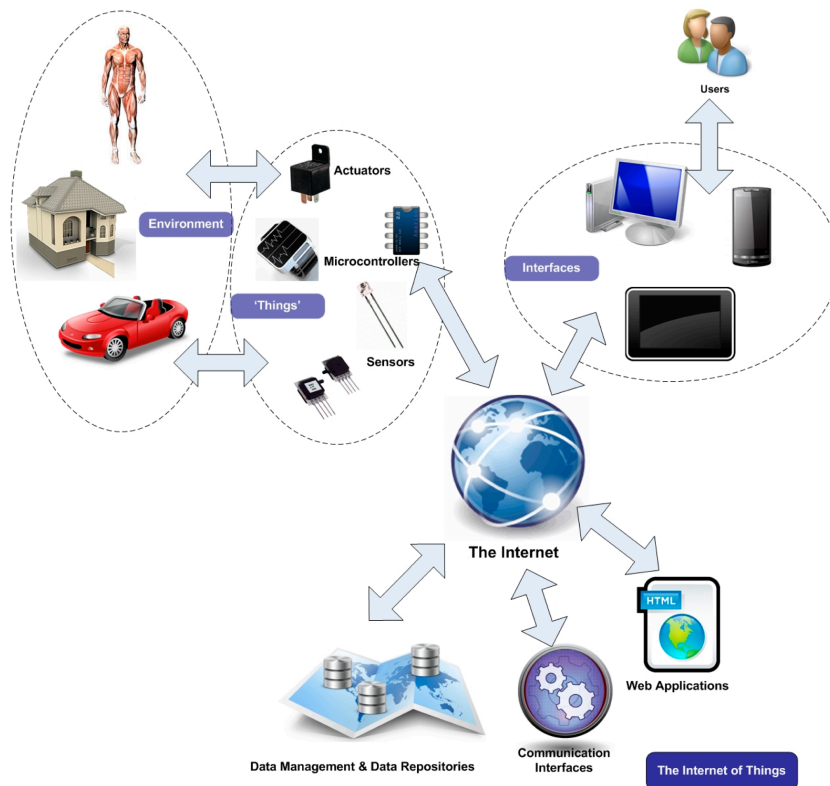
Design and Implement a Scalable, Real-time Data Processing Platform: Develop an architecture using AWS services that can dynamically scale to process high volumes of IoT data in real-time, accommodating fluctuations in data volume and velocity.

Achieve Low-latency Data Processing and Analysis: Ensure the platform can process and analyze data with minimal latency, enabling timely insights and actions for IoT applications.

Facilitate Seamless Data Integration and Management: Create a flexible data processing pipeline capable of handling diverse data formats and integrating data from various IoT sources, simplifying data management challenges.

Optimize Cost-efficiency: Implement cost-optimization strategies within the platform to manage the expenses associated with cloud resources, ensuring the solution remains cost-effective at scale.

Ensure Data Security and Privacy: Incorporate robust security measures and privacy practices to protect data throughout the processing and analytics pipeline, adhering to compliance standards and best practices.



Literature Review

1. ["ClouT: Leveraging Cloud Computing Techniques for Improving Management of Massive IoT Data"](#) - This paper discusses the synergy between cloud computing and IoT data analytics, highlighting the necessity for scalable, flexible architectures to process and analyze IoT data streams in real time.
2. ["Real-time Data Processing in IoT Networks: Challenges and Opportunities"](#) - Focuses on the specific challenges of real-time data processing within IoT networks, such as latency, data heterogeneity, and scalability. The paper reviews current solutions and identifies opportunities for improvement, particularly through the use of cloud-based serverless computing models.
3. ["Serverless computing for Internet of Things: A systematic literature review"](#) - Provides a comparison of different serverless architectures for IoT data processing, evaluating their performance, scalability, and cost-effectiveness. It emphasizes the benefits of AWS services in creating efficient data pipelines for real-time analytics.
4. ["Enabling privacy and security in Cloud of Things: Architecture, applications, security & privacy challenges"](#) - Addresses the critical aspects of security and privacy in IoT data analytics within the cloud. This paper outlines best practices for securing data pipelines, from ingestion to analytics, and the role of encryption, access control, and compliance standards.
5. ["Real-time analytics on IoT data"](#) - Investigates the transformative potential of real-time data analytics in IoT, providing case studies across different industries. It underscores the role of cloud platforms in enabling dynamic data analytics capabilities.

Progress Report -

Methodology

Phase 1: Design and Architecture

System Architecture Design: Define a scalable, modular architecture incorporating AWS IoT Core for device connectivity, Amazon Kinesis for real-time data streaming, AWS Lambda for serverless data processing, and Amazon Redshift or DynamoDB for analytics and storage. **Data Flow Mapping:** Outline the complete data journey, from ingestion from IoT devices, through transformation and processing, to analytics and storage, ensuring a clear path for data throughout the system.

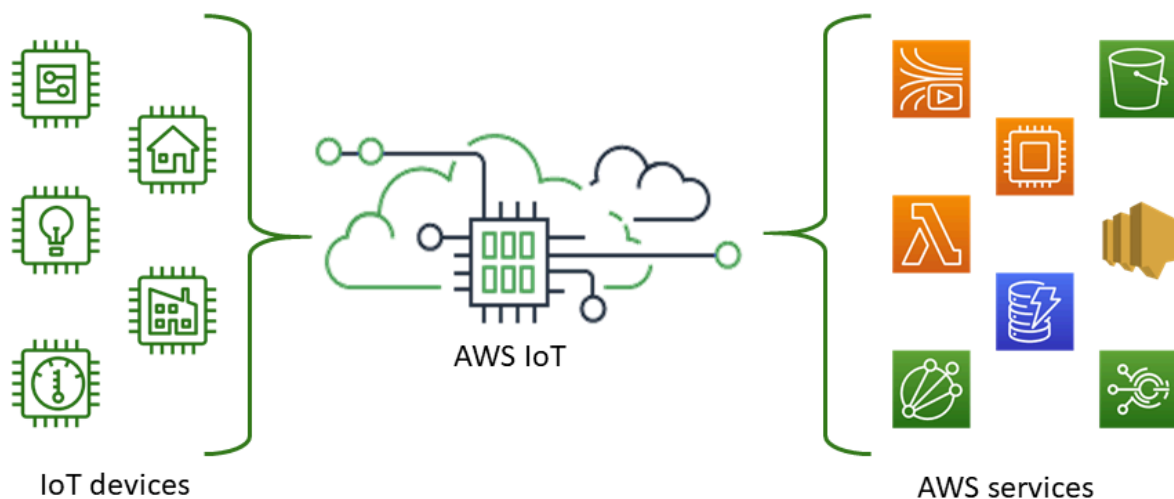
Phase 2: Implementation

AWS Environment Setup: Configure the necessary AWS services, ensuring they are securely integrated and optimized for cost and performance.

Data Ingestion and Streaming: Implement the ingestion process using **AWS IoT Core**, set up data streams with **Amazon Kinesis**, and ensure the system can handle data at the expected velocity and volume.

Serverless Processing Functions: Develop **AWS Lambda** functions to process and transform data as it flows through the pipeline, implementing any necessary business logic or data cleansing operations.

Analytics and Storage Configuration: Set up **Amazon Redshift or DynamoDB** for storing and analyzing processed data, designing schemas and analytics queries tailored to the project's goals.



Phase 3: Testing, Evaluation, and Optimization

Functional Testing: Verify that each component of the system works as intended, from data ingestion to analytics.

Performance and Scalability Testing: Conduct load testing to evaluate the system's performance under varying data volumes and processing demands, adjusting configurations as needed for optimal scalability.

Cost Evaluation: Monitor and analyze the costs associated with running the platform, identifying areas for optimization to ensure cost-efficiency.

IoT Device Consideration:

For an AWS-based real-time data analytics project, numerous IoT devices tailored to specific industries can enhance data collection and analysis:

Environmental Monitoring: Utilize weather stations for climate data and air quality sensors for pollution levels, aiding research and smart city applications.

Healthcare: Employ wearable health monitors for vital signs and smart inhalers for medication tracking, supporting personal health and patient care.

Smart Homes: Implement smart thermostats for energy efficiency and smart locks/cameras for enhanced security.

Industrial IoT: Deploy predictive maintenance sensors for machinery upkeep and asset tracking systems for logistics optimization.

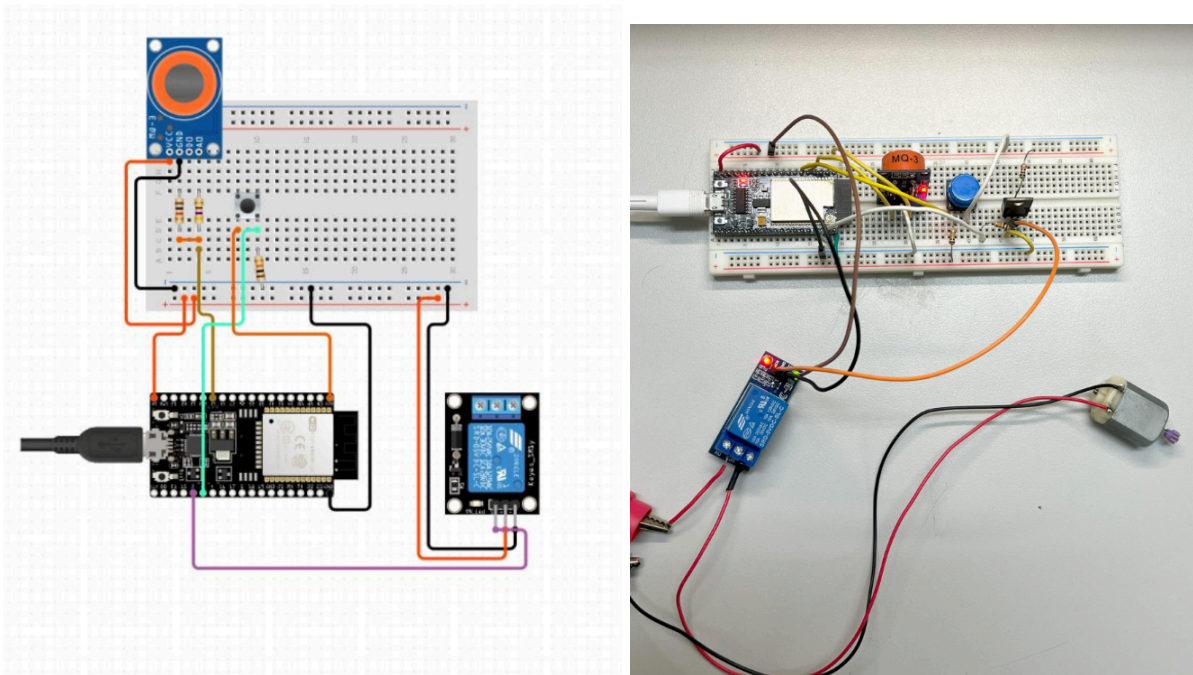
Agriculture: Use soil moisture sensors for irrigation management and livestock tracking devices for herd health.

Retail: Integrate smart shelves for inventory management and footfall sensors for customer behavior insights.

Smart Cities: Install waste management sensors for efficient collection and parking sensors to reduce congestion.

- Ensure IoT devices support required connectivity protocols (Wi-Fi, Bluetooth, LoRaWAN) for reliable data transmission.
- Confirm compatibility with AWS IoT Core and other AWS services.
- Plan for diverse data volumes and types, ensuring robust data handling and analysis capabilities.
- Prioritize security to protect sensitive information transmitted by IoT devices.

Considerable IoT Device: : Last year, we did a project on ‘**An IoT based Alcohol Sensing System For Safe Driving**’. The system employs an ESP32 Microcontroller integrated with an MQ3 Alcohol Sensor and a Relay module to detect alcohol levels in breath. Programmed for synchronized functioning, it showcases results on the NETPIE Dashboard and sends notifications. If alcohol levels exceed a predefined threshold, it prevents the vehicle's engine from starting. Additionally, a push button is incorporated to activate the sensor or reset the system. **For the real-time analytics project initially envisioned to use NETPIE for dashboarding and notifications, transitioning to AWS services can enhance scalability, security, and integration capabilities, particularly for IoT applications.**



System Diagram of IoT based Alcohol Detection System


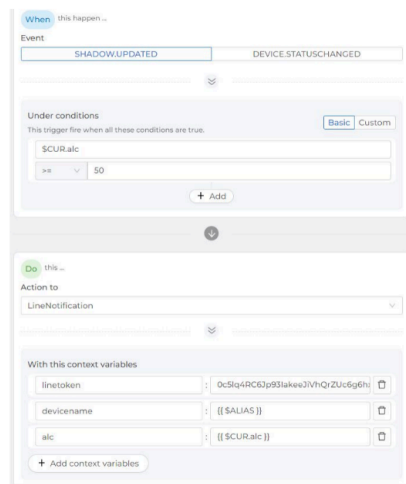

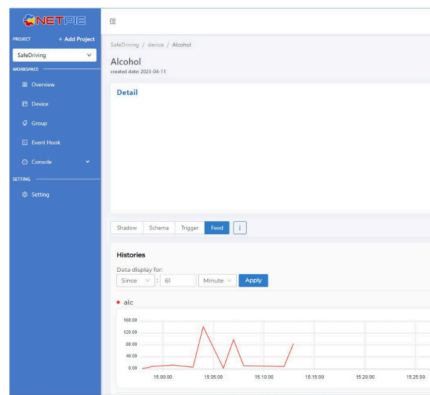
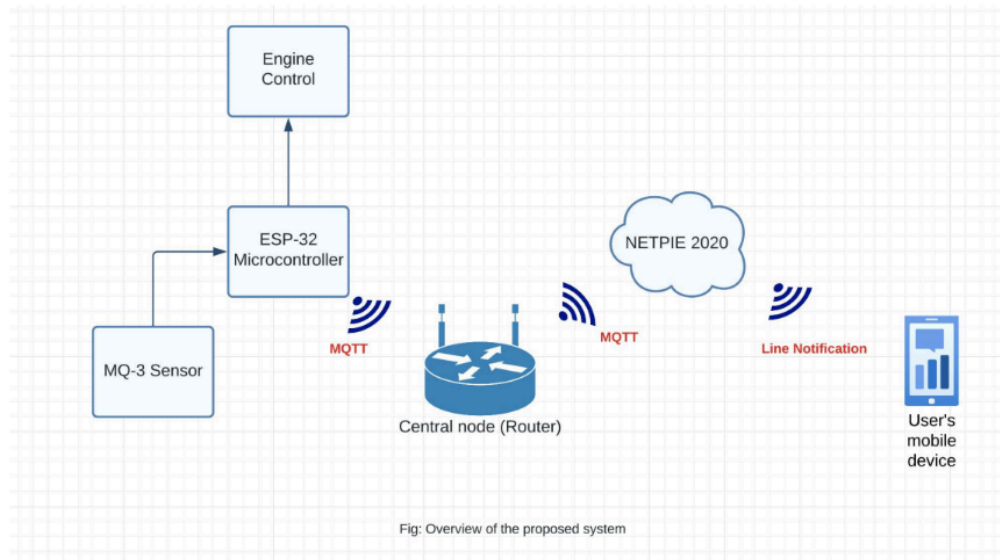


Figure - The old system architecture within the NETPIE interface.

Transition to AWS:

AWS IoT Core:

- **Device Integration:** ESP32 connects to AWS IoT Core, enabling secure communication between the IoT device and AWS cloud services.
- **MQTT Protocol:** Utilizes MQTT messaging protocol for lightweight, real-time communication of sensor data.

Amazon Kinesis Data Streams:

- **Real-time Data Processing:** Streams alcohol sensor data in real-time, facilitating immediate analysis and decision-making processes.

AWS Lambda:

- **Event-driven Logic:** Processes incoming data from Amazon Kinesis, executing logic based on alcohol level detection (e.g., comparing against threshold levels).
- **Control Commands:** Sends commands to activate the relay module to prevent engine start-up if necessary, through AWS IoT Core.

Amazon DynamoDB:

- **Data Storage:** Stores historical data of alcohol level readings for further analysis or auditing purposes.

Amazon SNS (Simple Notification Service):

- **Alerts and Notifications:** Sends real-time alerts or notifications to relevant stakeholders if alcohol levels exceed the predefined threshold.

Amazon QuickSight:

- **Data Visualization and Dashboarding:** Integrates with DynamoDB for visualizing historical data and real-time analytics on a user-friendly dashboard.

System Operation:

- When the push button is pressed, the ESP32 activates the MQ3 Alcohol Sensor.
- If alcohol is detected above the threshold level, the system prevents the vehicle's engine from starting by not activating the relay module.
- Sends real-time data to AWS IoT Core, which is processed and analyzed through AWS Lambda and Amazon Kinesis.
- Triggers an Amazon SNS notification to a predefined contact list.
- Logs the event in Amazon DynamoDB for record-keeping.
- Data analytics and visualization on Amazon QuickSight allow for monitoring and analyzing trends over time.

Benefits of using AWS compared to Netpie.

In the exploration of cloud platforms suitable for IoT applications, an analysis was conducted comparing AWS (Amazon Web Services) IoT Core and NetPIE. This comparison sought to understand the capabilities, scalability, security measures, and integration potential of each platform, highlighting the areas where AWS IoT Core often surpasses NetPIE, particularly for projects with extensive scalability and complex integration needs. The findings are presented from a passive assignment submission perspective, emphasizing the inherent strengths of AWS IoT Core in various dimensions.

Scalability and Reliability: The investigation revealed that AWS IoT Core is inherently capable of supporting a significant scale of operations, accommodating billions of devices and trillions of messages effortlessly. Its infrastructure is designed for high availability and redundancy, crucial for the rapid scaling of IoT applications. In contrast, while NetPIE is adept at handling many applications, its scalability does not parallel the enterprise-level capabilities offered by AWS.

Integration with Cloud Services and Ecosystem: AWS IoT Core's seamless integration with a broad array of AWS services stands out as a significant advantage. This integration facilitates the development of complex IoT solutions that can leverage advanced analytics, machine learning, and storage services directly within the AWS ecosystem. On the other hand, NetPIE, although it offers integration capabilities, lacks the extensive ecosystem provided by AWS.

Security: The analysis highlighted the robust security features of AWS IoT Core, including mutual authentication, fine-grained authorization, and encryption, both at rest and in transit. The comprehensive and continuously audited security model of AWS IoT Core provides a level of security assurance that is paramount for IoT projects, especially those handling sensitive data.

Developer Tools and Community Support: The vast AWS developer community and the plethora of documentation, tutorials, forums, and professional support options available were identified as significant benefits. AWS offers IoT-specific developer tools that facilitate device management, simulation, and testing, advantages that are less pronounced with NetPIE due to its smaller scale community and support resources.

Global Infrastructure: For projects targeting a global audience, AWS IoT Core's extensive data center presence around the world offers a strategic advantage, enabling localized data processing and adherence to regional data privacy regulations. This global infrastructure is a critical factor for applications requiring low latency and compliance with local data laws, areas where NetPIE may not provide the same level of support.

Cost and Pricing Model: While AWS IoT Core operates on a pay-as-you-go pricing model, offering cost-effectiveness at scale, careful management of resources is required to optimize expenses. The transparency in AWS pricing aids in effective planning and cost management, an aspect that may be more straightforward with NetPIE's pricing structure, particularly for smaller or initial-stage projects.

Current Status of the Project:

Design and Architecture Planning: A scalable system architecture incorporating AWS IoT Core for device connectivity was outlined, laying the groundwork for secure and efficient IoT device integration and data processing.

AWS Environment Setup: AWS services required for the project were configured, ensuring secure integration and optimization for both cost and performance, facilitating the seamless operation of the analytics platform.

Data Ingestion and Streaming Implementation: The system was enabled to ingest data from IoT devices, notably an ESP32 microcontroller with an MQ3 Alcohol Sensor, through AWS IoT Core, setting the stage for real-time data processing and analysis.

Integration of IoT Devices with AWS Services: An ESP32 microcontroller, equipped with an MQ3 Alcohol Sensor, was successfully integrated with AWS IoT Core. This critical step ensures secure, real-time communication between the IoT device and AWS cloud services using the MQTT protocol.

Initial Code Implementation and Testing: A foundational codebase was developed for the ESP32, enabling it to collect data from the MQ3 sensor, assess alcohol concentration, and accordingly control a servo motor. Additionally, this code supports connecting to AWS IoT Core and publishing the collected sensor data, marking a significant milestone in the project's development phase.

***Code has been shared in the provided github link.**

```
Alcohol now (mg/dL): 0.04      Status: Normal
Connecting to Wi-Fi
.....Connecting to AWS IoT
AWS IoT Connected!
Calibrating please wait..... done!.
Alcohol now (mg/dL): 0.04      Status: Normal
Alcohol now (mg/dL): 0.03      Status: Normal
Alcohol now (mg/dL): 0.04      Status: Normal
Alcohol now (mg/dL): 0.04      Status: Normal
Alcohol now (mg/dL): 0.04      Status: Normal
Alcohol now (mg/dL): 0.04      Status: Normal
```

The screenshot shows the AWS IoT console interface. At the top, there are two green success messages: "You successfully created thing ESP32_MQ3" and "You successfully created certificate dcae51132ee94f5835d2ef183a9f20912ed1fc67e86007e33e83eaf64d26c11a". Below these, the "Activity" tab is selected, showing a list of MQTT messages. The first message is a "Connected" event from the thing ESP32_MQ3, timestamped April 01, 2024, 00:41:58 (UTC+07:00). The second message is a "Disconnected" event from the same thing, timestamped April 01, 2024, 00:40:48 (UTC+07:00). The left sidebar contains navigation links for Test, Manage, and Device software. The top navigation bar shows the user's profile and the current region (N. Virginia).

us-east-1.console.aws.amazon.com/iot/home?region=us-east-14/thing/ESP32_MQ3

Test

- Device Advisor
- MQTT test client
- Device Location [New](#)

Manage

- All devices
 - Things
 - Thing groups
 - Thing types
 - Fleet metrics
 - Greengrass devices
 - LPWAN devices
 - Software packages [New](#)
 - Remote actions
 - Message routing
 - Retained messages
 - Security
 - Fleet Hub
- Device software
- Billing groups
- Settings
- Feature spotlight
- Documentation [🔗](#)

Tell us what you think

You successfully created thing ESP32_MQ3. [View thing](#)

You successfully created certificate dcae51132ee94f5835d2ef183a9f20912ed1fc67e86007e33e83eaf64d26c11a. [View certificate](#)

Attributes Certificates Thing groups Device Shadows **Activity** Packages and versions Jobs Alarms Defender metrics

Activity (5) [Info](#) [Clear](#) [MQTT test client](#)

Lists the most recent MQTT messages related to Device Shadow activity since you opened the thing details page. To see more messages related to this activity, choose the MQTT test client button.

▼ **Connected** April 01, 2024, 00:41:58 (UTC+07:00)

[\\$aws/events/presence/connected/ESP32_MQ3](#)

```
{
  "clientId": "ESP32_MQ3",
  "timestamp": 1711906918176,
  "eventType": "connected",
  "sessionId": "80ea07e9-dc20-457d-897c-7c1bfcc32798",
  "principalIdentifier": "dcae51132ee94f5835d2ef183a9f20912ed1fc67e86007e33e83eaf64d26c11a",
  "ipAddress": "40.228.96.198",
  "versionNumber": 4
}
```

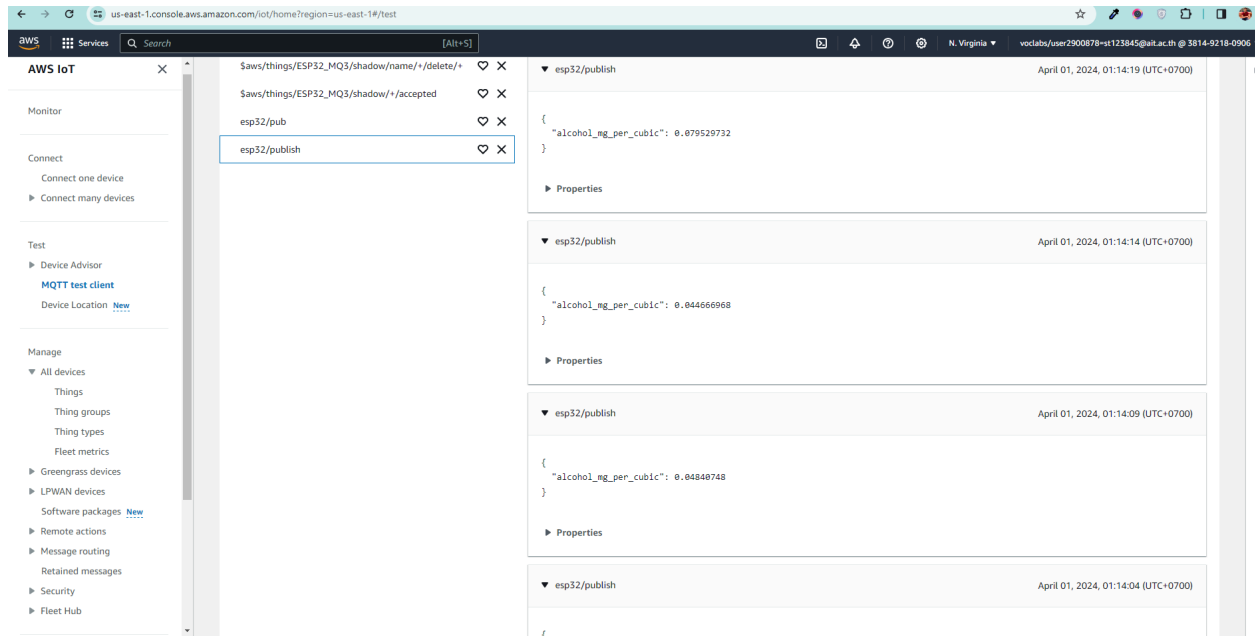
▼ **Disconnected** April 01, 2024, 00:40:48 (UTC+07:00)

[\\$aws/events/presence/disconnected/ESP32_MQ3](#)

```
{
  "clientId": "ESP32_MQ3",
  "timestamp": 1711906848021,

```

Successful Integration of ESP32 and Sensor with AWS IoT Core



Live Alcohol Level Data through MQTT Client

Ongoing Developments:

Integration with Amazon Kinesis and AWS Services: The incorporation of Amazon Kinesis for real-time data streaming, alongside AWS Lambda for serverless data processing and Amazon Redshift or DynamoDB for analytics and storage, is in progress. This critical step underpins the system's scalability and flexibility.

Serverless Processing Functions Development: Development of AWS Lambda functions to process and transform the data pipeline is advancing. These functions are being tailored to include logic for detecting alcohol levels and, if necessary, issuing commands to activate a relay module through AWS IoT Core.

Analytics and Storage Configuration: Configuration efforts for Amazon Redshift or DynamoDB are ongoing, with a focus on designing schemas and analytics queries. This will facilitate the extraction of actionable insights from the processed IoT data.

Real-time Data Processing Capabilities: The implementation of real-time data processing, utilizing Amazon Kinesis to stream alcohol sensor data, is being refined. This development is pivotal for supporting applications that necessitate rapid analysis and decision-making based on incoming data.

Data Visualization and Dashboarding: Plans for integrating Amazon QuickSight for data visualization and dashboarding are being finalized. This integration aims to enhance monitoring

and analysis of data trends over time, providing an intuitive interface for accessing insights derived from the IoT data.

GPS Module Integration: Updates to the device firmware to enable reading of GPS coordinates are slated for completion, leveraging available libraries to streamline the process.

AWS Lambda Function for Notifications: A pivotal modification of the AWS Lambda function is underway to facilitate notification dispatch based on high alcohol level detections, incorporating device GPS location in the alerts. This requires extracting GPS data and alcohol level alerts from incoming messages, with AWS SNS or third-party APIs being configured to send notifications to pre-defined contacts, including potentially the local police department, subject to necessary consents and legal considerations.

Notification Process Upon Threshold Exceedance: A protocol is being established for the ESP32 to publish messages to AWS IoT Core, including alcohol levels and GPS coordinates, upon threshold exceedance. This will trigger the modified Lambda function, ensuring timely notifications that include critical location information.

Alignment with AWS Well-Architected Framework

Operational Excellence

- The project incorporates detailed planning, testing, and optimization phases, ensuring secure and efficient AWS service configuration.
- Emphasis on continuous improvement is evident through functional, performance, and scalability testing, alongside cost optimization efforts.

Security

- Security measures are prioritized, especially concerning sensitive IoT device data management, utilizing AWS IoT Core and MQTT protocol for secure device connectivity and messaging.
- AWS Lambda and Amazon DynamoDB are employed for secure data processing and storage, adhering to best security practices.

Reliability

- Utilizes AWS IoT Core and Amazon Kinesis for robust real-time data streaming and processing capabilities, essential for handling large volumes of IoT data.
- The serverless architecture with AWS Lambda ensures system adaptability to demand changes without manual intervention, enhancing reliability.

Performance Efficiency

- The project leverages serverless computing, real-time data processing with Amazon Kinesis, and efficient data storage and analytics with DynamoDB or Amazon Redshift.
- The architecture's design for minimal latency in data processing showcases a commitment to maintaining high performance efficiency.

Cost Optimization

- A proactive approach to managing cloud resource costs is highlighted, focusing on identifying and implementing optimization strategies to maintain cost-efficiency at scale.
- The adoption of serverless computing contributes to cost optimization by automatically adjusting resource allocation based on demand.

Sustainability

- The project's efficient use of cloud resources and serverless computing indirectly supports sustainability by optimizing resource usage and reducing environmental impacts.

Cost Estimation:

AWS Pricing Calculator

>

My Estimate

My Estimate

Edit

Export

Share

Estimate date: April 1, 2024

Estimate summary

Info

Upfront cost

0.00 USD

Monthly cost

155.15 USD

Total 12 months cost

1,861.80 USD

Includes upfront cost

Getting Started with AWS

Get started for free

Contact Sales

My Estimate

Duplicate

Delete

Move to

Create group

Add support

Add service

Find resources

<input type="checkbox"/>	Service Name	Status	Upfront cost	Monthly cost	Description	Region	Config Summary
<input type="checkbox"/>	AWS IoT Core	-	0.00 USD	111.01 USD	-	Asia Pacific (Singapore)	Number of devices (MQTT) (2)...
<input type="checkbox"/>	AWS Lambda	-	0.00 USD	7.80 USD	-	Asia Pacific (Singapore)	Architecture (x86), Architectur...
<input type="checkbox"/>	Amazon Kinesis Data Streams	-	0.00 USD	0.00 USD	-	Asia Pacific (Singapore)	Duration of data retention (1 ...
<input type="checkbox"/>	Amazon QuickSight	-	0.00 USD	31.40 USD	-	Asia Pacific (Singapore)	Number of working days per ...
<input type="checkbox"/>	Amazon Simple Notification S...	-	0.00 USD	4.94 USD	-	Asia Pacific (Singapore)	Requests (5 million per month...

Estimate Summary

- Upfront Cost: \$0 is indicated, suggesting services primarily based on a pay-as-you-go pricing model are being utilized, without any initial commitment costs.
- Monthly Cost: \$155.15 is seen, reflecting a manageable operational cost monthly for the scale of services being used.
- Total 12 Months Cost: \$1861.80 is provided, giving a clear view of the expected operational expenditure annually for budgeting and financial planning purposes.

Detailed Estimate Breakdown

MQTT: A substantial portion of the monthly cost, \$111.01, is attributed to the MQTT service. With a high volume of messages (10,000,000) and device registry requests (10,000,000), the optimization of message sizes and frequencies is suggested to ensure value is maximized from every interaction within the IoT solution.

AWS Lambda: A cost-effective solution for executing backend logic in response to events at scale is represented by \$7.8 monthly for 10,000,000 requests. Optimizing function execution time and considering concurrency requirements could lead to further cost reductions.

Amazon Kinesis Data Streams: No cost is associated with this service in the estimate, potentially due to the specific configuration (e.g., data retention of 1 day). If included in the architecture for real-time data streaming, the efficiency of stream processing and dynamic scaling with the workload is recommended for cost management.

Amazon QuickSight: A valuable tool for data-driven decisions is reflected by \$31.4 monthly, mainly for data visualization and analysis. Adjusting SPICE capacity based on actual usage and regularly reviewing the number of readers and authors could optimize these costs.

Standard Topics: A cost of \$4.94 monthly, likely related to Amazon SNS for notifications, is noted. Ensuring messages are targeted and necessary by regularly reviewing the notification strategy is recommended for cost management.

Acknowledgement

The note that AWS Pricing Calculator's estimates do not include taxes and depend on actual usage is crucial. It is suggested that actual usage is monitored through AWS Cost Management tools and compared with these estimates to manage the budget effectively.

Additional Costs incurred so far:

ESP32: 4 USD

MQ3 Sensor : 3 USD

SERVO Motor: 1 USD

Conclusion and Future Work

This project demonstrates the effective integration of IoT devices with AWS services to create a scalable, real-time data analytics platform, using an alcohol detection system as a prime example. The adaptability and efficiency of this approach highlight its applicability across various IoT applications, from environmental monitoring to smart homes and healthcare.

- **Expand IoT Device Range:** Explore the integration of additional IoT devices to diversify application fields.
- **Enhance Data Analytics:** Implement more sophisticated analytics algorithms for deeper insights.
- **Improve User Interaction:** Develop advanced user interfaces and dashboard features for better data visualization and interaction.
- **Explore Edge Computing:** Incorporate edge computing to reduce latency and network load for critical applications.
- **Strengthen Security Measures:** Continuously update security protocols to protect against evolving threats.

References:

1. **SECURITY AND PRIVACY FOR CLOUD AND IoT** - by Kuo-Hui Yeh, Weizhi Meng, SK Hafizul Islam, Kuan Zhang ,Ennan Zhai - University of Nebraska - Lincoln
2. **Real-time Data Processing Using AWS Lambda** by Tara E Walker
3. **A Comparative Study of Cloud-Connected IoT Solutions** by Durga Prasad Rai and Ramesh Mishra - Avadh University Ayodhya (U.P.)- India
4. **Alcohol Detection System in Vehicle Using Arduino** by Deepak Garg, Ayush Kumar Srivastava, Deep Paliwal, Shashank Shekhar, Ashutosh Singh Chauhan
5. **Accident Vehicle Automatic Detection System By Image Processing Technology ©IEEE 1994** - by A.ISuge, H.Takigawa, H.Osuga, H.Soma, K.Morisaki
6. Vehicle Navigation & Information Systems Conference Paul Baskett , Yi Shang , Michael V. Patterson Timothy Trull, **Towards A System for Body-Area Sensing and Detection of Alcohol Craving and Mood Dvs regulation © 2013 IEEE**